

Parallel & Distributed Computing: Lecture 35

Alberto Paoluzzi

January 7, 2020

- 1 Introduction to CUDA
- 2 Julia packages for GPU programming
- 3 CuArrays.jl


Introduction to CUDA

Hardware view

At the top-level:

- PCIe graphics card with a many-core GPU and
- high-speed graphics “device” memory

sits inside a standard PC/server with one or two multicore CPUs:



DGX-1 (Computational Sciences Lab – DMF)

The NVIDIA DGX-1 comes with a base operating system consisting of

- Ubuntu OS,
- Docker,
- Docker Engine Utility for NVIDIA GPUs,
- NVIDIA drivers.


This system is designed to run a number of NVIDIA-optimized deep learning framework applications packaged in Docker containers.

DGX-1 Ip: tesla2.dia.uniroma3.it

- 2 Intel® Xeon®

E5-2698 v4, 20-core, 2.2GHz, 135W

- 8 Tesla P100 (Pascal architecture)
 - 170 teraflops, FP16
 - 16 GB memory per GPU
 - 28,672 NVIDIA CUDA® Cores



building block is a “streaming multiprocessor” (SM):

- 192 cores and 64k registers
- 64 KB of shared memory / L1 cache
- 8 KB cache for constants
- 48 KB texture cache for read-only arrays
- up to 2K threads per SM

Multithreading

cores in a SM are SIMT (Single Instruction Multiple Threads) cores:

- groups of 32 cores (warp) execute the same instructions simultaneously, but with different data
- 32 threads all doing the same thing at the same time
- natural for graphics processing and much scientific computing

SIMT is a natural choice for many-core chips to simplify each core

Software view

The **master process** runs on the **CPU** and performs the following steps:

- 1 **initialises** card
- 2 **allocates memory** in host and on device
- 3 **copies data** from host to device memory
- 4 **launches multiple instances** of execution “**kernel**” on **device**
- 5 **copies data** from device memory to host
- 6 **repeats 3-5** as needed
- 7 **de-allocates all memory** and **terminates**

Software view

At a lower level, **within the GPU**:

- ① each instance of the **execution kernel** executes **on a SM**
- ② if # of **kernel instances** exceeds # SMs,
 - then **more than one instance** will run **on each SM** if there are enough registers and shared memory,
 - and the others **will wait** in a **queue** and **execute later**
- ③ **all threads** within **one instance** can access **local shared memory** but can't see what the other instances are doing (even if they are on the same SM)
- ④ there are **no guarantees** on **the order** in which **the instances execute**

CUDA programming

CUDA (Compute Unified Device Architecture) is NVIDIA's program development environment:

- based on C/C++ with some extensions
- FORTRAN support provided by compiler from PGI (owned by NVIDIA) and by IBM XL compiler
- lots of example code and good documentation
- large user community on NVIDIA forums

CUDA advantages

CUDA has several advantages over traditional general-purpose computation on GPUs (GPGPU) using graphics APIs:

- Scattered reads – code can read from arbitrary addresses in memory
- Unified virtual memory (CUDA 4.0 and above)
- Unified memory (CUDA 6.0 and above)
- Shared memory – CUDA exposes a fast shared memory region that can be shared among threads. This can be used as a user-managed cache, enabling higher bandwidth than is possible using texture lookups.
- Faster downloads and readbacks to and from the GPU
- Full support for integer and bitwise operations, including integer texture lookups

Cuda programming

Installing CUDA on a system, there are 3 components:

- driver
 - low-level software that controls the graphics card
- toolkit
 - nvcc CUDA compiler
 - Nsight IDE plugin for Eclipse or Visual Studio profiling and debugging tools
 - several libraries
- SDK
 - lots of demonstration examples
 - some error-checking utilities
 - not officially supported by NVIDIA almost no documentation

Cuda programming (from <https://en.wikipedia.org/wiki/CUDA>)

CUDA 8.0 comes with the following libraries (in alphabetical order):

- **cuBLAS** – CUDA Basic Linear Algebra Subroutines library, see main and docs
- **CUDART** – CUDA Runtime library, see docs
- **cuFFT** – CUDA Fast Fourier Transform library, see main and docs
- **cuRAND** – CUDA Random Number Generation library, see main and docs
- **cuSOLVER** – CUDA based collection of dense and sparse direct solvers, see main and docs
- **cuSPARSE** – CUDA Sparse Matrix library, see main and docs
- **NPP** – NVIDIA Performance Primitives library, see main and docs
- **nvGRAPH** – NVIDIA Graph Analytics library, see main and docs
- **NVML** – NVIDIA Management Library, see main and docs
- **NVRTC** – NVIDIA Runtime Compilation library for CUDA C++, see docs

Tutorial 01: Say Hello to CUDA (C language)

🏠 CUDA Tutorial

Search docs

Home

Tutorials

Tutorial 01: Say Hello to CUDA

Tutorial 01: Say Hello to CUDA

Introduction

A quick comparison between
CUDA and C

Compiling CUDA programs

Putting things in actions.

Exercise: Converting vector
addition to CUDA

Profiling performance

Wrap up

Acknowledgments

Tutorial 02: CUDA in Actions



YOUR AD HERE

[< Previous](#)
[Next >](#)
[Docs](#) » [Tutorials](#) » Tutorial 01: Say Hello to CUDA

Tutorial 01: Say Hello to CUDA

Introduction

This tutorial is an introduction for writing your first CUDA C program and offload computation to a GPU. We will use CUDA runtime API throughout this tutorial.

CUDA is a platform and programming model for CUDA-enabled GPUs. The platform exposes GPUs for general purpose computing. CUDA provides C/C++ language extension and APIs for programming and managing GPUs.

In CUDA programming, both CPUs and GPUs are used for computing. Typically, we refer to CPU and GPU system as *host* and *device*, respectively. CPUs and GPUs are separated platforms with their own memory space. Typically, we run serial workload on CPU and offload parallel computation to GPUs.

A quick comparison between CUDA and C

Following table compares a hello world program in C and CUDA side-by-side.

C

```
void c_hello(){
    printf("Hello World!\n");
}

int main() {
    c_hello();
    return 0;
}
```

CUDA

```
__global__ void cuda_hello(){
    printf("Hello World from GPU!\n");
}

int main() {
    cuda_hello<<<1,1>>>();
    return 0;
}
```

[v: latest](#)

Julia packages for GPU programming

<https://devblogs.nvidia.com/gpu-computing-julia-programming-language/>

Julia is already well regarded for programming **multicore CPUs** and **large parallel computing** systems,

recent developments **make the language suited** for **GPU computing** as well.

The **performance** possibilities of GPUs can be **democratized** by providing more high-level tools that are **easy to use** by a large community of **applied mathematicians** and **machine learning programmers**.

In this blog post, **I will focus** on **native GPU programming** with a Julia package that enhances the **Julia compiler** with **native PTX code** generation capabilities: **CUDAnative.jl**

High-Performance GPU Computing in the Julia Programming Language

JuliaGPU — GPU Computing in Julia



JuliaGPU

GPU Computing in Julia

<https://juliagpu.org/>

Verified

Repositories 28

Packages

People 8

Projects

Pinned repositories

 **CuArrays.jl**

A Curious Cumulation of CUDA Cuisine

Julia ★ 214 🍴 74

 **CUDAnative.jl**

Julia support for native CUDA programming

Julia ★ 344 🍴 50

 **OpenCL.jl**

OpenCL Julia bindings

Julia ★ 184 🍴 32

 **AMDGPUnative.jl**

Julia interface to AMD/Radeon GPUs

Julia ★ 28 🍴 1

 **ArrayFire.jl**

Julia wrapper for the ArrayFire library

Julia ★ 156 🍴 30

 **juliagpu.org**

The JuliaGPU landing page.

HTML ★ 3

JuliaGPU — CUDA programming in Julia

CUDA programming in Julia

This repository hosts a [Julia package](#) that [bundles functionality](#) from [several other packages](#) for CUDA programming, and provides [high-level documentation](#) and [tutorials](#) for effectively [using CUDA GPUs from Julia](#). The documentation is accessible at juliagpu.gitlab.io.

[CUDA.jl](#) includes functionality from the following packages:

- 1 [CUDAdrv.jl](#): interface to the CUDA driver
- 2 [CUDAnative.jl](#): kernel programming capabilities
- 3 [CuArrays.jl](#): GPU array abstraction

CuArrays.jl

[API reference documentation](#) of [CuArrays.jl](#). This documentation is a work in progress. For general usage instructions of CuArrays.jl and the rest of the Julia CUDA toolchain, please refer to the [CUDA.jl documentation](#).

CuArrays.jl

CuArrays.jl

CuArrays.jl provides a fully-functional GPU array, which can give significant speedups over normal arrays without code changes. CuArrays are implemented fully in Julia, making the implementation elegant and extremely generic

Read chapter 7 of book:

Avik Sengupta, [Julia High Performance](#): Optimizations, distributed computing, multithreading, and GPU programming with Julia 1.0 and beyond, 2nd Edition, Pakt>, 2019

Getting started with GPUs

CUDA, PTX and its use from Julia

CuArrays - Arrays on the GPU

Analyzing the performance of GPU code